

# Video Based Contextual Question Answering

Akash Ganesan  
akaberto@umich.edu

Divyansh Pal  
divpal@umich.edu

Karthik Muthuraman  
mkarthik@umich.edu

Shubham Dash  
shudbhamd@umich.edu

## CCS CONCEPTS

• **Computing methodologies** → **Object detection;**  
**Natural language generation;** **Scene understanding;**  
**Visual content-based indexing and retrieval;**

## 1 PROBLEM DEFINITION

The problem we are primarily looking to solve aims at building a contextual Question-Answering model for a given video. The current methodologies currently provides a robust model to handle Question-Answering in images, but we are trying to generalize this approach to be applied on videos. Adding on it the question-answering model, what we are trying to build should also be able to handle contextual queries or questions, for example if a frame has an image of a man and a cat sitting, the model should be able to handle queries like “where was the cat sitting with respect to the man” or “what is the man holding in his right hand”, questions which deal with the context of that particular frame in consideration.

The questions being asked will focus primarily on binary yes/no question answering as well as the questions which ask the “What”, “Why”, “How”, “When” and “Where”.

## 2 CHALLENGES

There are multiple challenges which we have to deal with in order to successfully build a model which can efficiently handle queries and answer those with respect to the video in consideration and also preserve the context embedded in the video. First challenge would be to build a contextual linking in between the scene graphs. Consider an example where a man is eating food in one frame which is considered as a key frame and in the second frame he is driving a car, so the link from one scene graph describing one frame and the second scene graph describing the other frame should be linked in such a way so that the graph when traversed for finding out the answers should be representative of how to video is evolving through time. Secondly, for gauging the efficiency of our proposed model we need to fix upon an evaluation metric which can provide a measure for calculating the efficiency of our model. One more interesting challenge which needs to be tackled for building a model which handles contextual question-answering on videos in general, we need to look at how scalable the model would be, for example how efficient would answer retrieval be if the video is a very long.

## 3 PRIOR WORK

Previous related work can be found in the related fields of video summarization [6], image captioning [1, 2] and scene

graph generation [3, 5]. We reviewed papers that propose methods to generate key-frames of interest from a long video. A major part of our project will draw from works relating to dense captioning of images and generating scene graphs. There is a wide body of recent literature which propose novel and optimized techniques for the aforementioned tasks. Most of these rely on Deep Learning methods for object detection and captioning and ML and optimization techniques to generate the scene graphs. Finally, in terms of evaluation processes, possible test datasets and possible metrics, there is a wide range of papers and open datasets that can be relevant to our project [4].

## 4 PROPOSED METHODOLOGY

In our methodology, we first use the video data and find key frames. For each keyframe, we do semantic segmentation to get different localized objects, which will serve as our nodes during graph generation. YOLO and Faster R-CNN are usually used for semantic segmentation for their speed and accuracy. For our scope we are considering a dataset that gives us semantic segmentation results on the frames of the video. We then use a dense captioning algorithm to generate captions for each frame based on the dense captioning algorithm [1]. Now, we can use generated captions to form a scene graph. These scene graphs from captions are generated using the algorithm mentioned in [5]. The next step is to link the scene graphs. This novel step will establish a relation between the existing scene graph and the incoming scene graph from the current frame. This will tell what has changed from one frame to the next and this relationship is important. We plan to do this by augmenting our generated scene graph with the difference of the new and the current scene graph.

We may ask questions on the image like, “When did the person get in the car”?, from two adjacent frames that has a person outside the car and one has the person inside it. So, the relationship between the frame needs to be captured. Individual questions like where is an object in a frame may be answered by scene analysis from questions. However, it is a challenge to ask for relationship between the frames that we plan to address.

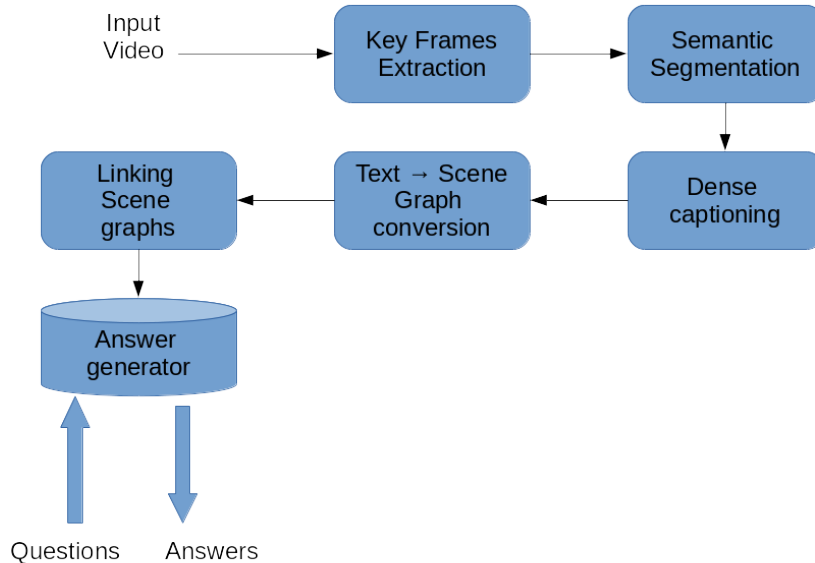


Figure 1: Model pipeline

## 5 DATASET DETAILS

The datasets which are considered for running our experiments are the Visual Genome and the Youtube-8M databases. The visual genome dataset has 108,077 images which has 75,729 unique objects, 40,480 unique objects and the number of unique relationships between those object, which form the nodes for the respective scene graphs as 40,513.

The second dataset which we will consider for testing our contextual question-answering model on videos is the Youtube-8M dataset. The dataset contains frame-by-frame annotations for eight million videos present.

## 6 EVALUATION CRITERIA

To evaluate the results, we will use Amazon’s mechanical Turk to generate answers to questions, answered by humans. This provides a baseline to check the accuracy of the generated answers from our question-answering model. There are several methods existing for evaluating the performance of image based question-answering models such as WUPS, which have been applied to video-based query answering models.

## 7 FUTURE WORK

Our work can be extended to incorporate speech content of video to generate more node edge combinations. This multimodal approach will make a denser graph but will store much more contextually rich information and can be used to answer much more in-depth questions. Once the graph is generated, a description text of the video can be generated. Other attributes of the object can be detected and incorporated to answer questions about emotion, expression, logic etc. Currently we focus mainly on actions and relationships but our work can be extended to emotion and inference based questions. Lastly, current video retrieval techniques rely heavily on video metadata such as video title/tags/description etc and less on the actual content/frames of the video. Extending our work, a video retrieval system can search on our representation of videos and hence the actual video content.

## REFERENCES

- [1] Justin Johnson, Andrej Karpathy, and Fei-Fei Li. 2015. DenseCap: Fully Convolutional Localization Networks for Dense Captioning. *CoRR* abs/1511.07571 (2015). arXiv:1511.07571 <http://arxiv.org/abs/1511.07571>
- [2] Andrej Karpathy and Fei-Fei Li. 2014. Deep Visual-Semantic Alignments for Generating Image Descriptions. *CoRR* abs/1412.2306 (2014). arXiv:1412.2306 <http://arxiv.org/abs/1412.2306>

- [3] Alejandro Newell and Jia Deng. 2017. Pixels to Graphs by Associative Embedding. *CoRR* abs/1706.07365 (2017). arXiv:1706.07365 <http://arxiv.org/abs/1706.07365>
- [4] Hyeonwoo Noh, Paul Hongsuck Seo, and Bohyung Han. 2016. Image Question Answering Using Convolutional Neural Network with Dynamic Parameter Prediction. In *2016 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2016, Las Vegas, NV, USA, June 27-30, 2016*. 30–38. <https://doi.org/10.1109/CVPR.2016.11>
- [5] Sebastian Schuster, Ranjay Krishna, Angel Chang, Li Fei-Fei, and Christopher D. Manning. 2015. Generating Semantically Precise Scene Graphs from Textual Descriptions for Improved Image Retrieval. In *Workshop on Vision and Language (VL15)*. Association for Computational Linguistics, Lisbon, Portugal.
- [6] Ke Zhang, Wei-Lun Chao, Fei Sha, and Kristen Grauman. 2016. Video Summarization with Long Short-term Memory. *CoRR* abs/1605.08110 (2016). arXiv:1605.08110 <http://arxiv.org/abs/1605.08110>